

American cities profile based on Foursquare data

Pedro Pessoa

Abstract

Here I present a study of the business profile of cities in the Americas. These profiles were constructed with the “Scrapping City Data Toolbox” (SCDT) python library to gather data for all Foursquare venues in each studied city. The Foursquare venue category is used as the type of business. We apply principal component analysis (PCA) in order to determine the most statistically relevant characteristics that differentiate the studied cities’ business ecosystems. The present manuscript can be seen as either (i) An explanation of SCDT with an example of its use or (ii) a preliminary study on the economic geography of cities in the Americas, for which the data was obtained using SCDT.

Keywords: Foursquare, Machine Learning, Principal component analysis, Mahalanobis.

Contents

1	Introduction	2
2	Data Acquisition	3
2.1	Creating the Grid	3
2.2	Obtaining the data in each city	4
3	Methods	7
3.1	Descriptive Statistics	8
3.2	Dimensional Reduction	8
3.3	Mahalanobis distance	10
4	Results	11
4.1	Dimensional reduction	11
4.2	Mahalanobis distance	11
5	Conclusion	14

1 Introduction

Foursquare is a location-based social network on which businesses (or venues, in the Foursquare nomenclature) can be registered and users can find their venues when close to their location, as well as write and read reviews left by other users. The searches are obtained by the Foursquare places application programming interface (FSAPI) [1] which can be called by any registered developer on free accounts. FSAPI, as a tool, is used by providing a search point (in geographical coordinates) and a search radius returning up to a hundred Foursquare venues withing a distance defined by the search radius of that search point.

The main motivation for the study presented here is to check how the data provided by FSAPI can give enough information about the “business ecosystem” of a city. For the scope of the present article “business ecosystem” means a list of every type of business – which here we will take as synonymous to the Foursquare venue category – and the fraction of businesses in that city that are of that type. However FSAPI is not programmed to give all venues in a city. As a way to adapt to this, I present the “Scrapping City Data Toolbox” (SCDT) [2], a python library developed by me that is able to, based on the geographical limits of a city, use FSAPI to give all Foursquare venues within the city.

For the study presented here I have chosen the ten most populated cities in the Americas – São Paulo, Ciudad de México, Lima, New York City, Bogotá, Rio de Janeiro , Santiago, Los Angeles, Caracas, and Buenos Aires – as our target cities, meaning the data for which our machine learning algorithms will be trained on. Then I chose five other cities as testing cities, meaning the ones for which the machine learning algorithms will compare the data to – namely Salvador, Toronto, Chicago, Guayaquil, and Medellín – these cities were chosen from the list of most populated cities in the Americas, following the ten target cities, so that no two of these cities are from the same country. Finally, Albany (NY, USA) was added as a testing city solely because is the city were the author of this report resides.

The data used in the present study was collected on Jan, 31 - Feb, 1 2021, and will not be made available in accordance to the Foursquare data retention policy [1]. However the GitHub repository for SCDT [2] contains all information necessary for one to reproduce the study done here.

I will apply the technique for dimensional reduction of data known as principal component analysis (PCA) and give graphical representations of the data for all studied cities in 2 PCA coordinates obtained from the target cities. Then I present other methods for differentiating the cities’ “business ecosystem” using that data, mainly calculating the Mahalanobis distances between the cities based on the target cities covariance matrix.

The layout of the paper is as follows: In the following section we will explore how the data was gathered, which is effectively a explanation of the algorithms presented in SCDT, first by selecting the set of points on which the FSAPI searches will be performed on, and how SCDT transforms the FSAPI data obtained from these searches into a single pandas dataframe. In section 3 we will describe the two methods used to differentiate the cities, mainly we will give a brief explanation of PCA and Mahalanobis distances. In section 4 we will present the results obtained using the FSAPI data for the studies and the Methods presented in section 3. In section 5 we state some discussion of our results. In the last section we present our conclusions.




Figure 1: Representation of a square lattice of side $L = \sqrt{2}R_s$, with a circle of radius R_s around each lattice point. Note that any point in the interior of the lattice is also in one of the circles.




Figure 2: Graphical representation for the quantities d_ϕ and d_λ . They are related to the coordinates in Mercator coordinate system by (3).

2 Data Acquisition

2.1 Creating the Grid

In order to obtain the data for all venues in a city, within the limitations of FSAPI, one needs to do a set of queries that (i) call FSAPI a small number of times and (ii) within a given search radius, it cover the whole area of the city. In this section we will explain our algorithm executed by the function `make_grid` in SCDT. It takes as parameters: (i) the polygon¹ of a city and (ii) a search radius, r_s ; and it returns a set of points for which it is mathematically guaranteed that any venues in the city can be found within a distance r_s to one of these points. Although the function name indicates this set of points is a grid, it does not necessarily have these points in a lattice, as we will examine soon.

Our first step consists of making a square lattice of side L . If $L \leq \sqrt{2}r_s$, any point in the interior of the lattice is at a distance smaller or equal to r_s to a lattice point. A of the geometrical proof of this fact is sketched in Fig. 1. A search done this way may obtain the same venue more than once, this can be seen in Fig. 1 for venues found in the intersection of two search circles. This is not an issue, since we opt to minimize the number of searches instead of the amount of data collected. However, for that reason, when treating the data one has to eliminate repeated venues.

The second step is a response to the fact that FSAPI receive the search parameters in terms of latitude, ϕ , and longitude, λ , this notation follows the Universal transversal Mercator coordinate system (see [4]). Its is important to understand that, in these coordinates, the space is no longer Euclidean. The distance d_ϕ between two points in the same longitude, namely (ϕ, λ) and $(\phi + \Delta\phi, \lambda)$, and the distance d_λ between two points in the same longitude, namely (ϕ, λ) and $(\phi, \lambda + \Delta\lambda)$, as represented in Fig.2 is (see [4])

¹By polygon we mean the set of limiting coordinates such as the ones obtained by Nominatim [3].

$$d_\phi = R_E \Delta\phi \quad \text{and} \quad d_\lambda = R_E \cos(\phi) \Delta\lambda, \quad (1)$$

where R_E is the radius of the Earth. Equation (3) assumes that $\Delta\phi$ and $\Delta\lambda$ are in natural units for angles (radians) also we do not consider changes in altitude as they are negligible compared to R_E .

The way to tackle this issues will be defining a rectangular lattice in Mercator coordinates with sides given by:

$$\Delta\phi = \frac{\sqrt{2}r_s}{R_E} \quad \text{and} \quad \Delta\lambda = \frac{\sqrt{2}r_s}{R_E} \frac{1}{c_{\max}} \quad (2)$$

where c_{\max} is the cosine for the latitude of the closest point to the equator in the city, mathematically defined a $c_{\max} \doteq \max\{\cos(\phi) | \phi \in P\}$ where P is the city's polygon if chosen this way, the sides of the lattice can be derived by substituting (2) into (3) obtaining

$$d_\phi = \sqrt{2}r_s \quad \text{and} \quad d_\lambda = \sqrt{2}r_s \frac{\cos(\phi)}{c_{\max}} \leq \sqrt{2}r_s, \quad (3)$$

as explained before, a lattice side smaller than $\sqrt{2}r_s$ is guaranteed to search all venues.

The third step is to remove the points in the lattice guaranteed not to give points within the city. In this, `make_grid` calls `grid_condition`, another function in SCDT that receives a point and the city's polygon and return the truth value if the point is either inside the city or immediate neighbour in the lattice to a point in the city and a false value otherwise. In Fig. 3 we represent how the points are removed in the example of São Paulo, in green are the points inside the city and in red are immediately outside, both are kept by `make_grid`, while the points in grey are removed. This guarantees all the area of the city will be searched, but also include some points up to $2r_s$ away from the city's boundary. Since the present work focus on big cities, this should not make a considerable difference and venues so close might still be considered part of the city business "ecosystem".

Having explained how the points for search are selected, we will proceed to explain how to use SCDT to obtain the venue data for a entire city.

2.2 Obtaining the data in each city

The SCDT function that gives the full data for a city is `gather_fsdata`. This is the only function that needs to be called once the set of points is done. The parameters of `gather_fsdata` are a set of points (as the ones generated by `make_grid` explained in the previous subsection), the search radius r_s (which should be the same as the one for which `make_grid` was called, although a bigger one can also be used with some loss of performance) and the users Foursquare credentials. The returned value of `gather_fsdata` is a pandas dataframe with the name, venue category and coordinates of each foursquare venue found in the city. Since there are daily limits to the usage of FSAPI, `gather_fsdata` estimates an upper quota of calls (which is the number of search points) and asks the user if they want to proceed, so that the user knows how many FSAPI calls will be made by SCDT.

For each of the points `gather_fsdata` calls the SCDT `collect_data` which takes the point's coordinate, search radius and credentials and returning the data for all FSAPI venues within the search radius of that point as a pandas data frame. This is done by a simple call of FSAPI at that point and appending the data with the columns representing (i) the venue's name, (ii) venue's category name (which roughly speaking is the type of business that venue is registred as) and (iii) the venue's coordinates. The conversion of the FSAPI into pandas data frame with these columns is done by the SCDT function `clear_data`. After doing so

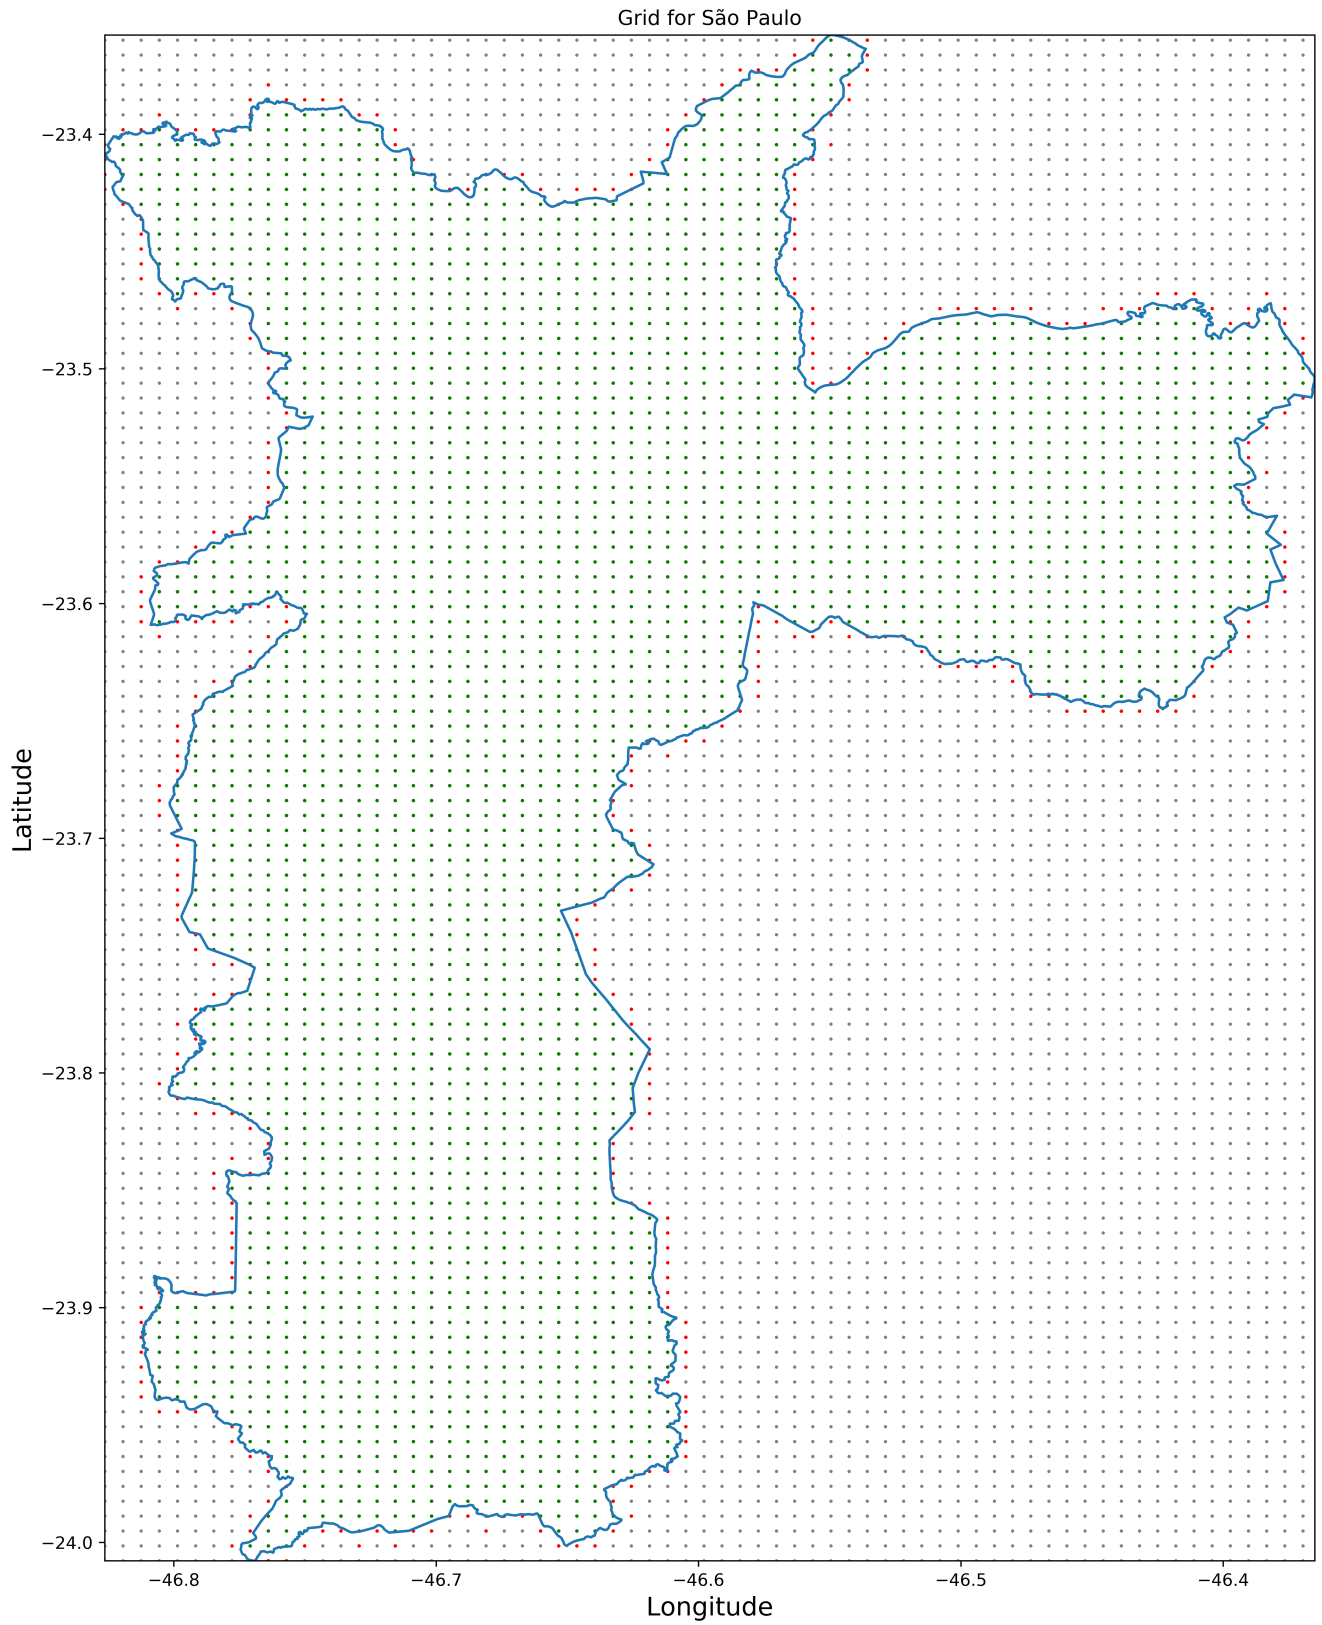


Figure 3: Representation for for São Paulo's map on how `make_grid` works. The points for which `grid_condition` is false (in grey) are removed.




Figure 4: Graphic representation for the second fail-safe in `gather_fsdata`. The search in the central circle in the Right-hand side turns is substituted by the search in the four blue circles in the left side.

for every single point, `gather_fsdata` condenses the data frame for each point into a single data frame that removes repetitions, which is returned.

The code for `gather_fsdata` has two fail-safe clauses: The first triggers when the call of FSAPI made by `collect_Data` fails. This may happen, for example, if the system loses internet connection or if the user exceeds their FSAPI daily limit². In that case, the fail-safe consists of calling the SCDT function `recerrorfix` that recursively tries to call `collect_Data` for the same point until it succeeds, at each attempt `recerrorfix` asks for a keyboard command so that the user may know (and fix) the what is causing the issue.

The second fail-safe is triggered when the limit of venues in a single FSAPI call is reached. This is a response to the issue that FSAPI has built-in a limit to the number of venues found in a single query. If a search returns a quantity of venues equal to the limit (100 venues within the current version of FSAPI), that implies that some venues in that area were not accounted for. Then `gather_fsdata` will cover all of that area with some searches on a smaller radius. This is done by creating a set of four points at a distance $r'_s = \frac{r_s}{\sqrt{2}}$ of the original point in the latitude-longitude grid. That means, around a point (ϕ, λ) `gather_fsdata` create a set of points p with coordinates

$$p = \{(\phi - \Delta\phi', \lambda), (\phi + \Delta\phi', \lambda), (\phi, \lambda - \Delta\lambda'), (\phi, \lambda + \Delta\lambda')\}, \quad (4)$$

where

$$\Delta\phi' = \frac{r'_s}{R_E} \quad \text{and} \quad \Delta\lambda = \frac{r'_s}{R_E} \frac{1}{\cos(\phi)}, \quad (5)$$

and calling `gather_fsdata` for this set of points and a new search radius r'_s .

The procedure for the second fail-safe is illustrated in Fig. 4 where one can see that all the area around the original point will be searched within r'_s of the points in p . The area is smaller, therefore fewer venues will return. Since `gather_fsdata` is called recursively this

²This is unlikely to be an issue. In the São Paulo example in section 2.1 for a $r_s = 500\text{m}$ there is 3427 points, which is much smaller number than the 99500 daily calls allowed by FSAPI for a free account.

	Coffee Shop	Bakery	Pizza Place	Gym Venue	Park
Target cities					
São Paulo	0.0340	0.0513	0.0459	0.0504	0.0084
Ciudad de México	0.0533	0.0271	0.0233	0.0343	0.0167
Lima	0.0357	0.0199	0.0199	0.0196	0.0587
New York City	0.0412	0.0202	0.0394	0.0230	0.0220
Bogotá	0.0518	0.0352	0.0279	0.0240	0.0285
Rio de Janeiro	0.0213	0.0347	0.0330	0.0532	0.0061
Santiago	0.0426	0.0309	0.0219	0.0218	0.0265
Los Angeles	0.0479	0.0129	0.0242	0.0211	0.0196
Caracas	0.0326	0.0685	0.0257	0.0283	0.0188
Buenos Aires	0.1025	0.0331	0.0496	0.0307	0.0083
Testing cities					
Salvador	0.0225	0.0392	0.0393	0.0438	0.0030
Toronto	0.0888	0.0191	0.0305	0.0198	0.0418
Chicago	0.0340	0.0133	0.0290	0.0206	0.0312
Guayaquil	0.0455	0.0148	0.0225	0.0237	0.0213
Medellín	0.0542	0.0214	0.0238	0.0290	0.0209
Albany	0.0269	0.0087	0.0400	0.0215	0.0237

Table 1: Coordinates for a few common venues categories in all studied cities. For the very first entry it means that 3.25% of the venues found in São Paulo are coffee shops.

way, even if the search done around one of the smaller circles also reaches the limit of venues the fail-safe triggers again. So that it does searches in even smaller search radius until it is possible to have fewer than the limit of venues in every search. With all the data from each of the smaller searches, we can guarantee that, despite the FSAPI limit, `gather_fsdata` will obtain the data for every venue in the city.

This completes the explanation on how SCDT can gather the data for the whole city, examples of usage can be seen in the notebooks found in GitHub repository for SCDT [2]. Now, with the data gathered from SCDT for the target cities, we can apply machine learning to them. This will be done in the following section.

3 Methods

Our goal is to use the data acquired through the method described in the previous section to give information on the “business ecosystem” of each city. The data for each city is converted as a proportion of each venue category. Meaning, for each city, there is a data point in a space for which each feature (coordinate) represent a category and the value is the fraction of the city’s venues of that are of the category. As an example, the features for a few common venue categories are presented on Table 1. Further work in cleaning the data was necessary before applying the machine learning techniques to the data set.

First, the category venue for restaurants that are of the same gentilic of the country in which the city resides are rewritten as ‘Restaurant’ (e.g. A venue of the category ‘Mexican Restaurant’ will be accounted as just ‘Restaurant’ in Ciudad de México, while still accounted for as ‘Mexican Restaurant’ in New York). Second, we gathered somewhat redundant venue categories to be summed as a single category, as implemented by the python dictionary

`simplify_dic` (e.g. venues of the Foursquare categories ‘Gym’ and ‘Gym / Fitness Center’ will be accounted as ‘Gym Venue’). Third, we remove the columns for venue categories that either (i) is only present in a single city³, or (ii) refers to an administrative city division that is not informative of the “Business ecosystem” (e.g. ‘Border Crossing’ or ‘Neighborhood’). Also, for a valid comparison, the data frames for the testing cities will only include the columns referring to venue categories present in the target cities.

3.1 Descriptive Statistics

We initiate the data analysis by obtaining a few descriptive statistics for our data set. First we present the five most common venue categories for each city on Table 2, second we present the correlations in the target cities between a few common venue categories on Table 3. A few preliminary results can be drawn from this. Bakeries is a common venue category in all target Latin American cities – except Lima – and do not appear as a common category for any Anglo-American cities. Also, pizza places are strongly positively correlated with gyms (cities with above average number of pizza places are likely to also have an above average quantity of gyms), while parks are strongly negatively correlated with gyms (cities with above average number parks are likely to also have a below average quantity of gyms).

3.2 Dimensional Reduction

When applying machine learning algorithm, some issues arise from the high dimensionality of the data. If we have n – in our case ten, one for each target city – data points on a space of d dimensions – more than 500 in our case – deters some machine learning methods - e.g. the covariance and correlation matrix has dimension d but its range cannot be larger than n . Hence, we need a method that gives the data in fewer dimensions, or a smaller number of features, but still better preserving the information for target cities. This is achieved by the technique known as principal component analysis (PCA), in this subsection we will outline the PCA used here (for a more general and in-depth explanation see [5]).

The dataframe will be represented in the matrix D which has d columns – representing each of our coordinates for venue categories – and n lines – representing each target city – from which we construct the matrix for the centered data Dc , meaning

$$Dc_{i,j} \doteq D_{i,j} - \frac{1}{n} \sum_{l=1}^n D_{l,j} . \quad (6)$$

In order to apply PCA we first need to find the singular value decomposition (SVD) of Dc , that mean finding matrices $U_{(n,n)}$ ⁴, $S_{(n,n)}$, and $V_{(d,d)}$ so that

$$Dc = USV^t , \quad (7)$$

where V^t represents the transpose of V , S is a diagonal matrix, and both U and V are orthogonal matrices, meaning $(U^t)^{-1} = U$ and $(V^t)^{-1} = V$.

In SVD, the terms in the diagonal S are placed in decreasing order, meaning the first columns of S represent the most important dimensions for the data. PCA consists of using only a number k of dimensions, related to the k first columns of S , effectively using the dimensional reduced data matrix Dr_k defined as

$$Dr_k \doteq US_{[:, :k]} , \quad (8)$$

³This is done since a venue category for a single city will refer to a peculiarity that cannot be properly compared to other cities.

⁴Notation note: $A_{(i,j)}$ means that the matrix A has i lines and j columns.

	Most Common Venue	2nd Most Common Venue	3rd Most Common Venue	4th Most Common Venue	5th Most Common Venue
Target cities					
São Paulo	Restaurant	Bakery	Gym Venue	Pizza Place	Coffee Shop
Ciudad de México	Restaurant	Taco Place	Coffee Shop	Gym Venue	Bakery
Lima	Restaurant	Park	Seafood Restaurant	Coffee Shop	Chinese Restaurant
New York City	Coffee Shop	Pizza Place	Deli / Bodega	Chinese Restaurant	Gym Venue
Bogotá	Restaurant	Coffee Shop	Burger Joint	Fast Food Restaurant	Bakery
Rio de Janeiro	Restaurant	Gym Venue	Bar	Bakery	Pizza Place
Santiago	Coffee Shop	Plaza	Restaurant	Bus Station	Bakery
Los Angeles	Coffee Shop	Mexican Restaurant	Restaurant	Pizza Place	Fast Food Restaurant
Caracas	Bakery	Italian Restaurant	Pharmacy	Coffee Shop	Restaurant
Buenos Aires	Coffee Shop	Restaurant	Pizza Place	Ice Cream Shop	Bakery
Testing cities					
Salvador	Restaurant	Gym Venue	Pizza Place	Bakery	Food Truck
Toronto	Coffee Shop	Park	Pizza Place	Restaurant	Sandwich Place
Chicago	Mexican Restaurant	Coffee Shop	Park	Pizza Place	Bar
Guayaquil	Seafood Restaurant	Coffee Shop	Pharmacy	Restaurant	BBQ Joint
Medellín	Coffee Shop	Restaurant	Hotel	Burger Joint	Bar
Albany	Pizza Place	Restaurant	Coffee Shop	Convenience Store	Sandwich Place

Table 2: List of most common venue categories in each of the studied cities.

	Coffee Shop	Bakery	Pizza Place	Gym Venue	Park
Coffee Shop	1.000000	-0.185530	0.439791	-0.255923	-0.189618
Bakery	-0.185530	1.000000	0.209982	0.414556	-0.353419
Pizza Place	0.439791	0.209982	1.000000	0.480712	-0.615918
Gym Venue	-0.255923	0.414556	0.480712	1.000000	-0.668948
Park	-0.189618	-0.353419	-0.615918	-0.668948	1.000000

Table 3: Correlations for a few common venue categories. This only considers the correlation in the target cities.

as the data⁵. That means, each line of Dr_k represents a target city and the k columns found on that line are the dimensionally reduced data. For the remainder of this text k takes the value 2, leading to a representation of points in Cartesian coordinates that can be graphed into a planar graph.

The data for the testing cities will undergo a similar transformation. From (8) and (7) we can see that

$$Dc = Dr_n V^t \iff Dr_n = Dc V . \quad (9)$$

If we call the matrix for the data in testing cities tD , the dimensionally reduced data for the testing cities, tDr_k , appropriate for comparison to the target data is

$$tDr_k = tDc V_{[:,k]} , \quad \text{where} \quad tDc_{i,j} \doteq tD_{i,j} - \frac{1}{n} \sum_{l=1}^n D_{l,j} , \quad (10)$$

note that tDc is centered around the means for the target data instead of the testing, which is coherent to the idea that we are changing the testing data coordinated to the dimensionally reduced coordinates created from the target data.

3.3 Mahalanobis distance

With the data in PCA coordinates obtained in Dc and tDc we can define the matrices $X_{(1,k)}^\mu$, where μ is an enumeration of the studied cities, so that X^μ is the μ -th line of Dr_k , $X_{1,j}^\mu \doteq Dr_{k \mu,j}$, if μ refers to a target city; and X^μ the $(\mu - n)$ -th line of tDr_k , $X_{1,j}^\mu \doteq tDr_{k \mu-n,j}$, if μ refers to a testing city. From that we may define the inner product

$$\langle A|B \rangle = AC^{-1}(B)^t = \sum_{i=1}^n \sum_{j=1}^n A_{1,i} C_{i,j}^{-1} B_{1,j} , \quad (11)$$

where $A_{(1,k)}$ and $B_{(1,k)}$ – hence are in the same vector space as X^μ – and $C_{(k,k)}$ is the covariance matrix obtained from the target cities’ dimensionally reduced data DR_k . The Mahalanobis distance between the cities enumerated by μ and ν is defined as $d(\mu, \nu) \doteq \sqrt{\langle X^\mu - X^\nu | X^\mu - X^\nu \rangle}$ ⁶.

Here we defined the Mahalanobis distance for the dimensionally reduced coordinates. It could not be defined from the original coordinates since $n < d$ – meaning we have more data features than data points – hence the covariance matrix for the non reduced data has a range $\leq n$ and therefore it is not invertible.

Having our methods described, we can present the results obtained from them in the following section.

⁵Notation note: $A_{[i,:j]}$ is the matrix obtained from the top i lines and the first j columns of another matrix A .

⁶It is interesting to state that Mahalanobis distance is the Euclidean distance between the point in standardized coordinates [5] – centered and with a covariance matrix being the identity matrix.

	MRV	2nd MRV	3rd MRV	4th MRV	5th MRV
PC 1 +	Chinese Restaurant	Mexican Restaurant	Grocery Store	Sandwich Place	Italian Restaurant
PC 1 -	Coffee Shop	Gym Venue	Seafood Restaurant	Taco Place	Restaurant
PC 2 +	Bakery	Pizza Place	Coffee Shop	Gym Venue	Plaza
PC 2 -	Taco Place	Chinese Restaurant	Fried Chicken Joint	Seafood Restaurant	Park

Table 4: The most relevant venue categories (MRV) on the PCA coordinates obtained from target cities data.

4 Results

4.1 Dimensional reduction

Before presenting the dimensionally reduced data in the PCA coordinates, we can gather their meaning by noticing that, according to (9), the first k lines of V^t are what a unitary vector in the PCA coordinates transforming into in the original coordinates. That gives some insight into how PCA works for this dataset. In Table 4 we present venue categories associated to the largest and lowest numbers in the first two lines of V^t . Roughly speaking, from what is presented in Table 4, a city that has above (below) average fraction of Chinese or Mexican restaurants tends to be assigned with a positive (negative) value for the first PCA coordinate, analogously a city that has above (below) average fraction of coffee shops or gyms tends to be assigned with a negative (positive) value for the first PCA coordinate.

The set of points representing the studied cities, for the $k = 2$ principal components coordinates, is graphed in Fig. 5. We can see that the Anglo-American cities cluster together in the region where the first principal component coordinate (PC1) is high and the second principal component coordinate (PC2) is average. On the same graph, Brazilian cities cluster together in high values of PC2 and average values of PC1. Comparing these results to Table 4, we see that Anglo-American cities are characterized by a large fraction of Chinese restaurants, Mexican restaurants and sandwich places, while Brazilian are characterized by a larger fraction of bakeries, pizza places and gyms. These results concur with the most common venues presented in Table 2.

There is no clear pattern for the Hispanic-American cities, rather there is a broad divergence between them. Santiago is placed near the Anglo-American – which may be explained by Chile’s proximity to US-American economic standards – while Buenos Aires is placed near Brazilian cities – which can be explained by the strong interlink between Argentinian and Brazilian economies. Ciudad de México and Lima are considerably different from all other cities being characterized by low values of, respectively, PC1 and PC2. Comparing these results to Table 4, we can see that Ciudad de México has a large fraction of coffee shops, gyms and taco places; Lima’s position is justified compared to Table 2 by its large fraction of parks and seafood restaurants.

4.2 Mahalanobis distance

The Mahalanobis distances between all studied cities are presented in Table 5. All Brazilian cities are within a distance of a unit of São Paulo. Also, all Anglo-American are with a Mahalanobis distance of .5 from New York City, although Santiagois also within this distance. Ciudad de México and Lima are strongly separated with a distance higher than unit to every other city.

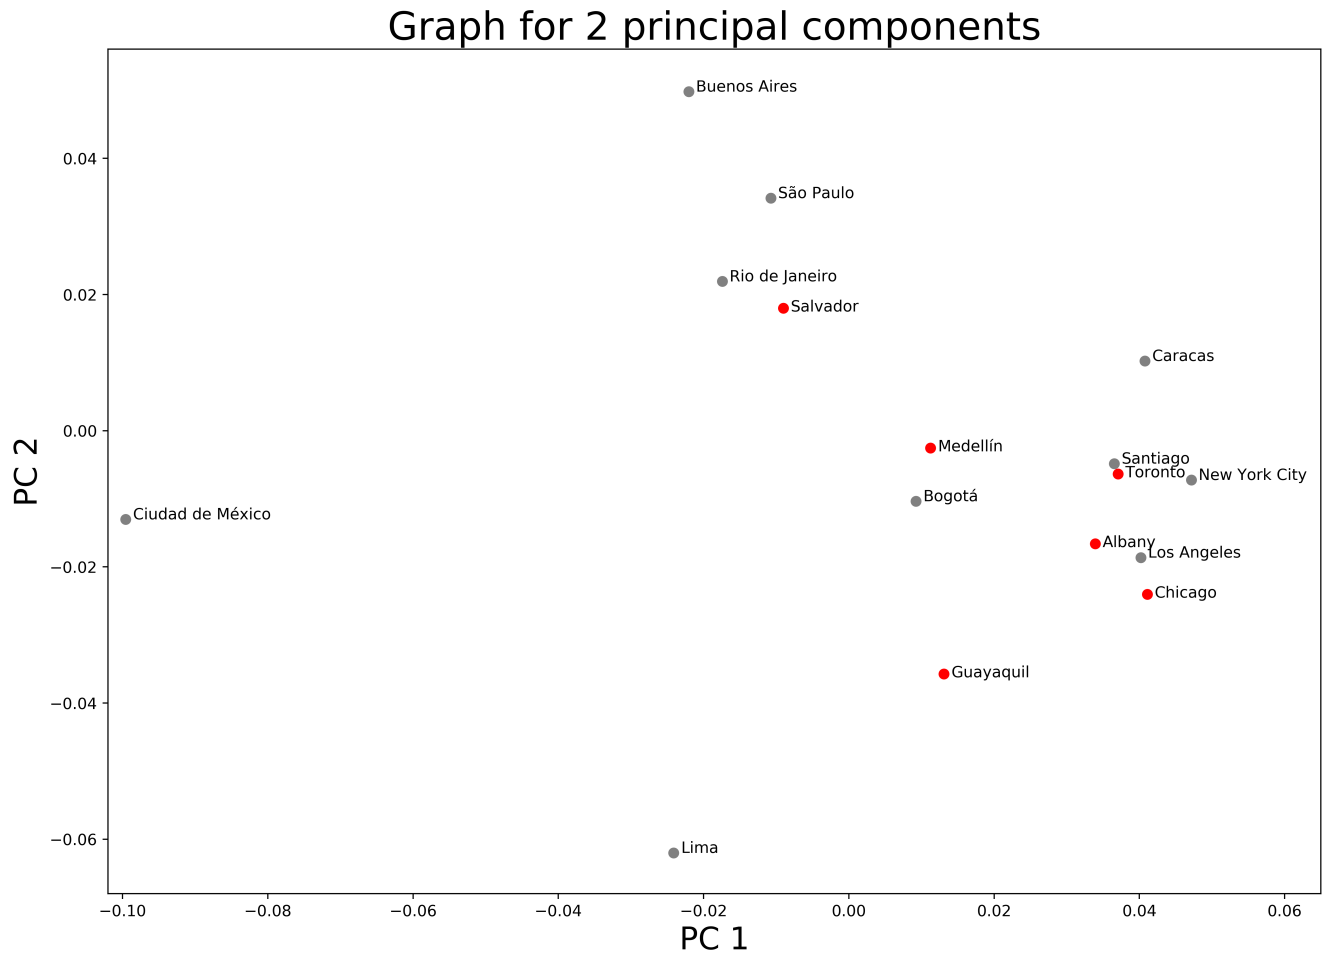


Figure 5: Cities data points in the PCA coordinates ($k = 2$) obtained for the data in the target cities. The target cities are represented in grey points while the testing cities are represented by red points.

	SP	CDMX	Lima	NYC	BGT	RIO	STG	LA	Caracas	BA	Salvador	Toronto	CCG	GYQ	Medellín	ALB
Target cities																
São Paulo	0.00	2.48	3.09	1.85	1.49	0.42	1.63	2.03	1.37	0.56	0.52	1.68	2.19	2.30	1.27	1.90
CDMX	2.48	0.00	2.29	3.26	2.41	2.14	3.03	3.10	3.20	2.65	2.24	3.04	3.14	2.60	2.48	2.96
Lima	3.09	2.29	0.00	2.36	1.81	2.69	2.27	1.99	2.72	3.58	2.58	2.24	1.89	1.18	2.06	1.94
New York City	1.85	3.26	2.36	0.00	0.85	1.71	0.25	0.40	0.58	2.38	1.48	0.23	0.55	1.18	0.81	0.42
Bogotá	1.49	2.41	1.81	0.85	0.00	1.19	0.63	0.74	0.96	2.05	0.99	0.63	0.83	0.82	0.25	0.58
Rio de Janeiro	0.42	2.14	2.69	1.71	1.19	0.00	1.47	1.82	1.34	0.90	0.23	1.51	1.96	1.97	1.01	1.68
Santiago	1.63	3.03	2.27	0.25	0.63	1.47	0.00	0.45	0.49	2.18	1.25	0.05	0.62	1.12	0.57	0.38
Los Angeles	2.03	3.10	1.99	0.40	0.74	1.82	0.45	0.00	0.92	2.59	1.60	0.40	0.17	0.81	0.82	0.15
Caracas	1.37	3.20	2.72	0.58	0.96	1.34	0.49	0.92	0.00	1.88	1.13	0.54	1.10	1.59	0.77	0.87
Buenos Aires	0.56	2.65	3.58	2.38	2.05	0.90	2.18	2.59	1.88	0.00	1.06	2.22	2.75	2.85	1.83	2.46
Testing cities																
Salvador	0.52	2.24	2.58	1.48	0.99	0.23	1.25	1.60	1.13	1.06	0.00	1.28	1.74	1.79	0.80	1.46
Toronto	1.68	3.04	2.24	0.23	0.63	1.51	0.05	0.40	0.54	2.22	1.28	0.00	0.57	1.08	0.59	0.34
Chicago	2.19	3.14	1.89	0.55	0.83	1.96	0.62	0.17	1.10	2.75	1.74	0.57	0.00	0.73	0.95	0.29
Guayaquil	2.30	2.60	1.18	1.18	0.82	1.97	1.12	0.81	1.59	2.85	1.79	1.08	0.73	0.00	1.06	0.77
Medellín	1.27	2.48	2.06	0.81	0.25	1.01	0.57	0.82	0.77	1.83	0.80	0.59	0.95	1.06	0.00	0.67
Albany	1.90	2.96	1.94	0.42	0.58	1.68	0.38	0.15	0.87	2.46	1.46	0.34	0.29	0.77	0.67	0.00

Table 5: Table of Mahalanobis distances between the studied cities based on the 2 PCA coordinates obtained from the target cities. The following abbreviations were used: São Paulo (SP), Ciudad de México (CDMX), New York City (NYC), Bogotá (BGT), Rio de Janeiro (RIO), Santiago (STG), Los Angeles (LA), Buenos Aires (BA), Chicago (CCG), Guayaquil (GYQ), and Albany (ALB)

5 Conclusion

Under the limitations provided by FSAPI, I was able to create SCDT – a python library that searches all Foursquare venues within a city. It works by turning the limits (or polygon) into a set of points so that FSAPI calls adequately covers the city, as explained in Section 2.1, and by gathering the data for all searches obtaining, effectively, data from a FSAPI search over the whole city, as explained in Section 2.2.

We transform the data into a table of fractions of the businesses in a city that are of the same type. This is posteriorly treated with PCA and the Mahalanobis distance between the cities in the two principal components is calculated. The meaning of the PCA coordinates can be visualized in Table 4. Fig 5 graphs the two principal components for all studied cities. There we can identify clusters of Brazilian and Anglo-American cities, Table 4, explains in terms of business types, how these agglomerations are formed. The Mahalanobis distance between the cities presented on Table 5 supports the conclusion.

References

- [1] Foursquare. Foursquare places API reference. <https://developer.foursquare.com/docs/places-api/>.
- [2] Pessoa, P. Scrapping city data toolbox - github repository. <https://github.com/PessoaP/ScrappingCityDataToolbox>.
- [3] Nominatim. Nominatim documentation. <https://nominatim.org/release-docs/develop/>.
- [4] Snyder, J. P. *Map projections: A working manual* (US Government Printing Office, 1987).
- [5] Barber, D. *Bayesian Reasoning and Machine Learning* (Cambridge University Press, 2012). Available at: <http://web4.cs.ucl.ac.uk/staff/D.Barber/textbook/200620.pdf>.